



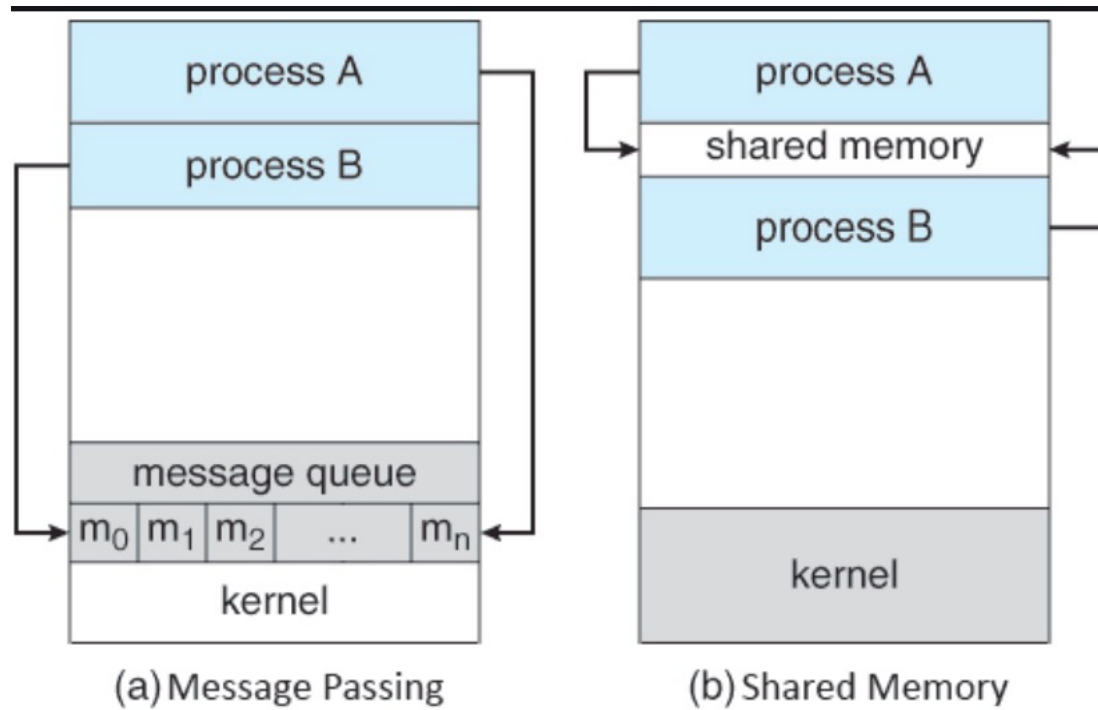
Chapter 6: Synchronization Tools

Background:

A cooperating process is one that can affect or be affected by other processes executing in the system.

- Cooperating processes can either directly share a logical address space (that is, both code and data) or be allowed to share data only through **shared memory** or **message passing**.
- Concurrent access to **shared data** may result in data **inconsistency**, however.
- Mechanisms to ensure the **orderly execution** of cooperating processes that share a logical address space.







Chapter 6-Objectives

- Critical-section problem
- Race conditions
- Hardware solutions to the critical-section problem
- Mutex locks, semaphores, monitors, and condition variables to solve the critical section problem






Cooperating processes

- CPU scheduler switches rapidly between processes to provide concurrent execution.
 - Processes may be interrupted **at any time**.
 - This means that one process may only partially complete execution before another process is scheduled.
 - Concurrent/parallel execution can contribute to issues involving the **integrity of data** shared by several processes
- Mechanisms to ensure the **orderly execution** of cooperating processes






Producer-consumer problem




- Producer process creates an item and adds it to the shared buffer.
- Consumer process takes items out of the shared buffer and “consumes” them.

- Suppose that we wanted to provide a solution to the consumer-producer problem with a limited buffer size:
 - We can do so by having an integer **counter** that keeps track of the buffer.
 - Initially, **counter** is set to 0.
 - It is **incremented by the producer** after it produces a new item and is **decremented by the consumer** after it consumes an item.









Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```


Circular buffer/queue






Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```







Race Condition

- **counter++** could be implemented as:

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** could be implemented as:

```
register2 = counter
register2 = register2 - 1
counter = register2
```

**outcome of the execution
depends on the particular
order in which the access
takes place!**

- Consider this execution **interleaving** with “count = 5” initially:

S0: producer execute	<code>register1 = counter</code>	{register1 = 5}
S1: producer execute	<code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute	<code>register2 = counter</code>	{register2 = 5}
S3: consumer execute	<code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute	<code>counter = register1</code>	{counter = 6}
S5: consumer execute	<code>counter = register2</code>	{counter = 4}


Two processes manipulate the variable counter concurrently!






Too much milk problem

time	You	Your Roommate
3:00	Arrive home	
3:05	Look in fridge, no milk	
3:10	Leave for grocery	
3:15		
3:20	Arrive at grocery	Arrive home
3:25	Buy milk	Look in fridge, no milk
3:35	Arrive home, put milk in fridge	Leave for grocery
3:45		Buy Milk
3:50		Arrive home, put up milk
3:50		Oh no!









“Withdraw at the same time from a joint account” Problem

Balance = \$300

ATM -1

\$300 ← Check Balance

??? ← Withdraw \$300

ATM-2

Check Balance → \$300

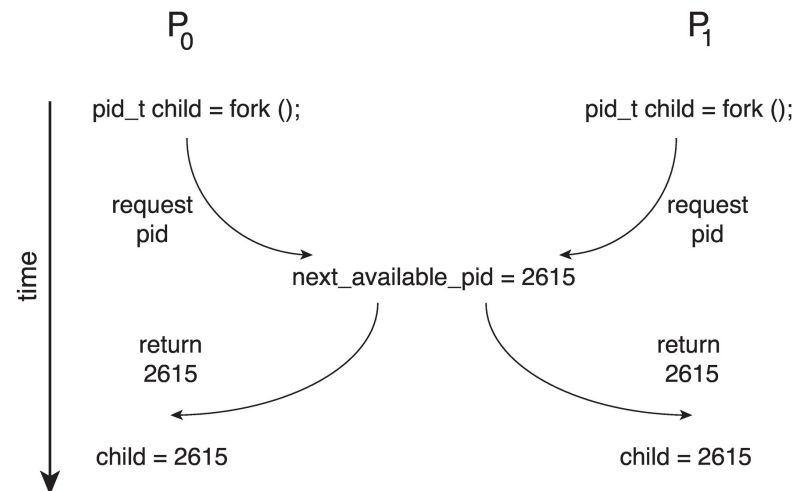
Withdraw \$300 → ???





Race Condition – kernel level example

- Processes P_0 and P_1 are creating child processes using the `fork()` system call.
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid).



- Unless there is **mutual exclusion**, the **same pid** could be assigned to two different processes!





Critical Section Problem

- Consider a system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$.
- Each process has **critical section** segment of code:
 - Process may be changing common/shared variables, updating table, writing file, etc
 - When one process in critical section, **no other process** may be in its critical section.
- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**.





Critical Section

General structure of process P_i

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```





Critical-Section Handling in OS

- The critical-section problem could be solved simply in a **single-core** environment if we could **prevent interrupts from occurring** while a shared variable was being modified.
- Unfortunately, this solution is **not feasible** in a **multiprocessor/multicore** environment.
 - Disabling interrupts on such systems can be **time consuming**, since this message must be passed to all the processors.





Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then **no other processes** can be executing in their critical sections.
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next **cannot be postponed indefinitely**.
3. **Bounded Waiting** - A **bound must exist** on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.
 - Assume that each process executes at a nonzero speed.
 - No assumption concerning **relative speed** of the n processes.






Critical-Section Handling in OS

Two approaches:


- **Preemptive kernel** – allows preemption of a process when running in kernel mode.
- It is possible for two kernel-mode processes to run simultaneously on different CPU cores.

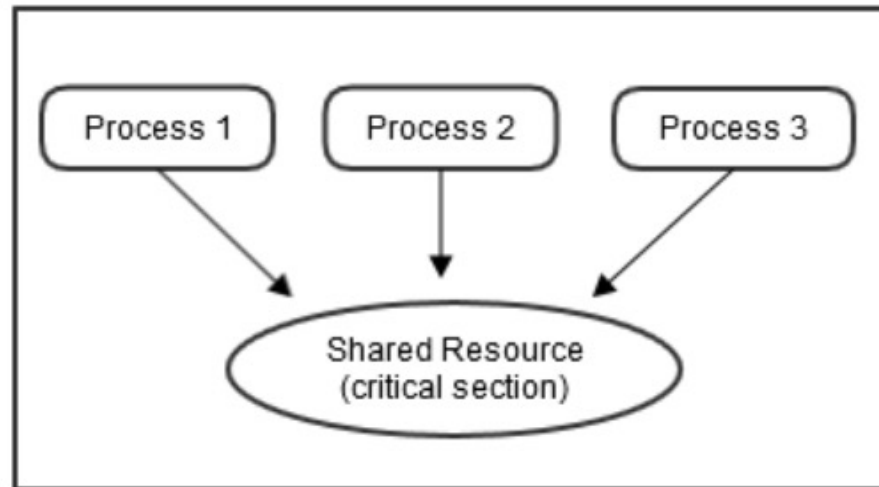


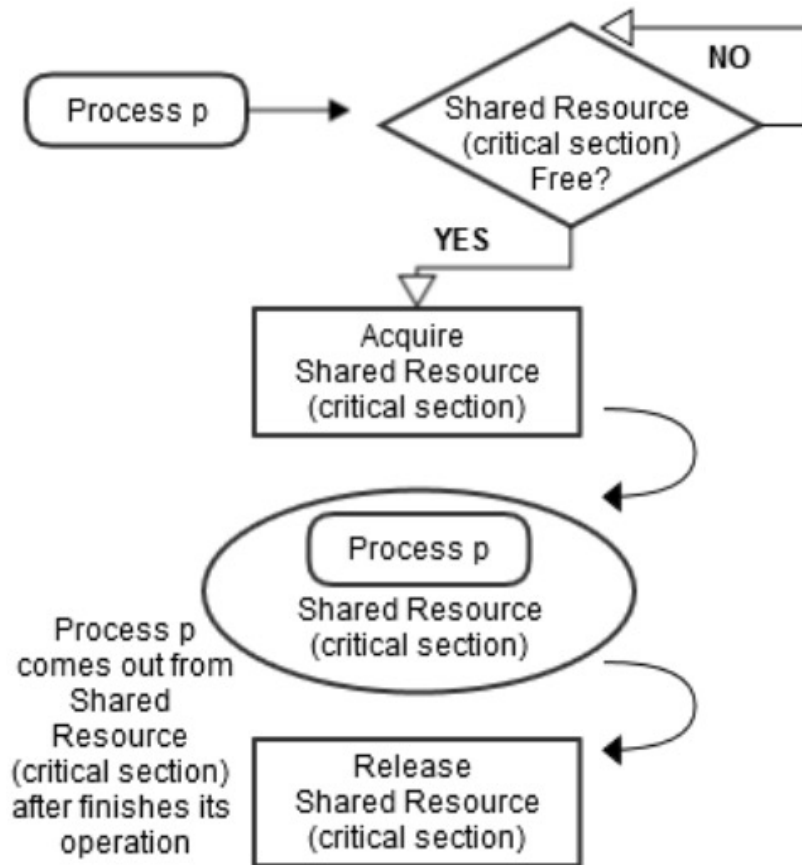


Critical-Section Handling in OS

- **Non-preemptive kernel** – process runs until exits kernel mode, blocks, or voluntarily yields the CPU
 - Essentially **free of race** conditions in kernel mode as only one process is active in the kernel at a time.
 - Non-preemptive kernels cannot be used for real-time scheduling or multiprocessor systems









Peterson's Solution

- Not guaranteed to work on modern architectures (but good algorithmic description of solving the problem)!
- Assume that the **load** and **store** machine-language instructions are atomic; that is, **cannot be interrupted**
- The two processes share two variables:
 - **int turn;**
 - **boolean flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i] = true** implies that process P_i is ready!





Algorithm for Process P_i

```
while (true){  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j)  
        ;  
  
    /* critical section */  
  
    flag[i] = false;  
  
    /* remainder section */  
  
}
```





Peterson's Solution (Cont.)

The three CS requirement are met:

1. Mutual exclusion is preserved

P_i enters CS only if:

either **flag[j] = false** or **turn = i**

2. Progress requirement is satisfied
3. Bounded-waiting requirement is met





Two-thread example

- Two threads share the data:

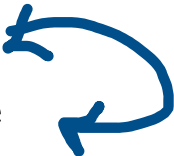
```
boolean flag = false;  
int x = 0;
```

- Thread 1 performs

```
while (!flag)  
    ;  
print x
```

- Thread 2 performs

```
x = 100;  
flag = true
```



- What is the expected output?



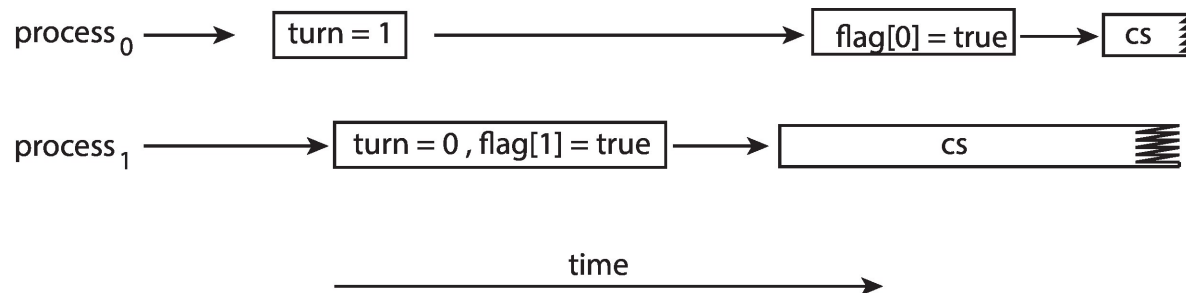


Peterson's Solution

Operations for Thread 2 may be **reordered**:

```
flag = true;  
x = 100;
```

- If this occurs, the output may be 0!
- The effects of instruction reordering in Peterson's Solution:



- This allows both processes to be in their critical section at the same time!
- *To improve system performance, processors and/or compilers may reorder read and write operations that have no dependencies.**





Synchronization Hardware

- Software-based solutions are **not guaranteed** to work on modern computer architectures.
- Many systems provide **hardware support** for implementing the critical section code.
 - These primitive operations can be used **directly** as synchronization tools, or they can be used to form the more abstract synchronization mechanisms.
- Three forms of hardware support:
 1. Memory barriers
 2. Hardware instructions
 3. Atomic variables





Memory Barriers

- Instruction that forces any change in memory to be propagated (made visible) to all other processors.
- When a memory barrier instruction is performed, the system ensures that all **loads and stores** are completed **before any subsequent load or store** operations are performed.





Memory Barrier

- We could add a memory barrier to the following instructions to ensure Thread 1 outputs 100 (the correct value):

- Thread 1:

```
while (!flag)
    memory_barrier();
print x
```

- Thread 2:

```
x = 100;
memory_barrier();
flag = true
```

* Low-level operations and are typically only used by kernel developers





Hardware Instructions

- Special hardware instructions that allow us to either **test-and-modify** the content of a word, or two *swap* the contents of two words **atomically** (uninterruptible)
- We can use these special instructions to solve the critical-section problem
- **Test-and-Set** instruction
- **Compare-and-Swap** instruction





test_and_set Instruction

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = true;
    return rv;
}
```

1. Executed **atomically**
2. Returns the **original value** of passed parameter
3. Set the **new value** of passed parameter to **true**

If two `test_and_set()` instructions are executed simultaneously (each on a different core), they will be executed sequentially in some arbitrary order.





Solution using test_and_set()

- Shared boolean variable **lock**, initialized to **false**
- Solution for a process P_i :

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
    /* remainder section */  
} while (true);
```





compare_and_swap Instruction

```
int compare_and_swap(int *value, int expected, int new_value)
{
    int temp = *value;

    if (*value == expected)
        *value = new_value;
    return temp;
}
```

1. Executed atomically
2. Returns the original value of passed parameter **value**
3. Set the variable **value** the value of the passed parameter **new_value** but only if ***value == expected** is true. That is, the swap takes place only under this condition.





Solution using compare_and_swap

- Shared integer `lock` initialized to 0;
- Solution:

```
while (true){  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

The first process that invokes `compare_and_swap()` will set `lock` to 1. It will then enter its critical section. Subsequent calls to `compare_and_swap()` will not succeed, because `lock` now is not equal to the expected value of 0.

Does not satisfy the bounded-waiting requirement

- Intel processors: `cmpxchg`





Atomic Variables

- **Atomic variable:** provides *atomic* (**uninterruptible**) updates on basic data types such as integers and booleans.
- For example, the **increment()** operation on the atomic variable **sequence** ensures **sequence** is incremented without interruption:

increment(&sequence) ;

- Atomic variables can be used to ensure **mutual exclusion** in situations where there may be a race on a single variable while it is being updated (e.g. a **counter** is incremented).





Atomic Variables

- The `increment()` function can be implemented as follows:

```
void increment(atomic_int *v)
{
    int temp;

    do {
        temp = *v;
    }
    while (temp != (compare_and_swap(v, temp, temp+1)) );
}
```

They do not entirely solve race conditions in all circumstances.





Mutex Locks

Previous solutions are **complicated** and **inaccessible** to application programmers. OS designers build **software tools** to solve **critical section** problem:

- Protect a critical section by first **acquire()** a lock then **release()** the lock
 - Boolean variable indicating if lock is available or not
- Calls to **acquire()** and **release()** must be **atomic**
 - Usually implemented via hardware atomic instructions such as compare-and-swap.
- **Busy waiting:** While a process is in its critical section, any other process that tries to enter its critical section must **loop continuously** in the call to `acquire()`.
- This lock therefore called a **spinlock**





Solution to Critical-section Problem Using Locks

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```





Mutex Lock Definitions

```
■ acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
  
■ release() {  
    available = true;  
}
```

These two functions must be implemented atomically.
Both `test-and-set` and `compare-and-swap` can be used to implement these functions.





Semaphore

Busy waiting: while a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the call to `acquire()`.

■ Semaphore S – integer variable: can only be accessed via **two indivisible (atomic)** operations

● `wait()` and `signal()`

Definition of the `wait()` operation:


```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

All modifications to the integer value of the semaphore in the `wait()` and `signal()` operations must be executed atomically

Definition of the `signal()` operation:

```
signal(S) {  
    S++;  
}
```







Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
- Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “**synch**” initialized to 0

P1 :

S_1 ;

signal (synch) ;


P1 and P2 share a common semaphore `synch`, initialized to 0.

P2 :

wait (synch) ;


S_2 ;









1. Process 1 is using the shared resource.



2. Process 2 is now in need of the shared resources.



3. Process 2 enters busy waiting till it has access to the shared resource.





Semaphore Implementation

- Rather than engaging in busy waiting, the process can **suspend** itself.
- Process is placed into a **waiting queue** associated with the semaphore, and the state of the process is switched to the waiting state. Then control is transferred to the CPU scheduler, which selects another process to execute.
- A process that is suspended, waiting on a semaphore S, should be **restarted** when some other process executes a `signal()` operation.
 - The process is restarted by a `wakeup()` operation, which changes the process from the waiting state to the ready state. The process is then placed in the ready queue.





Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- `typedef struct {
 int value;
 struct process *list;
} semaphore;`





Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```





Monitors

- Hide mutual exclusion: provides a convenient and effective mechanism for process synchronization
 - *Abstract data type*, internal variables only accessible by code within the procedure
 - Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    function P1 (...) { ... }

    function P2 (...) { ... }

    function Pn (...) {.....}

    initialization code (...) { ... }
}
```


A function defined within a monitor can access only those variables declared locally within the monitor and its formal parameters.

- Similarly, the local variables of a monitor can be accessed by only the local functions.





Schematic view of a Monitor



only one process at a time is active within the monitor





Condition Variables

- Monitors: insufficient for modeling some synchronization schemes

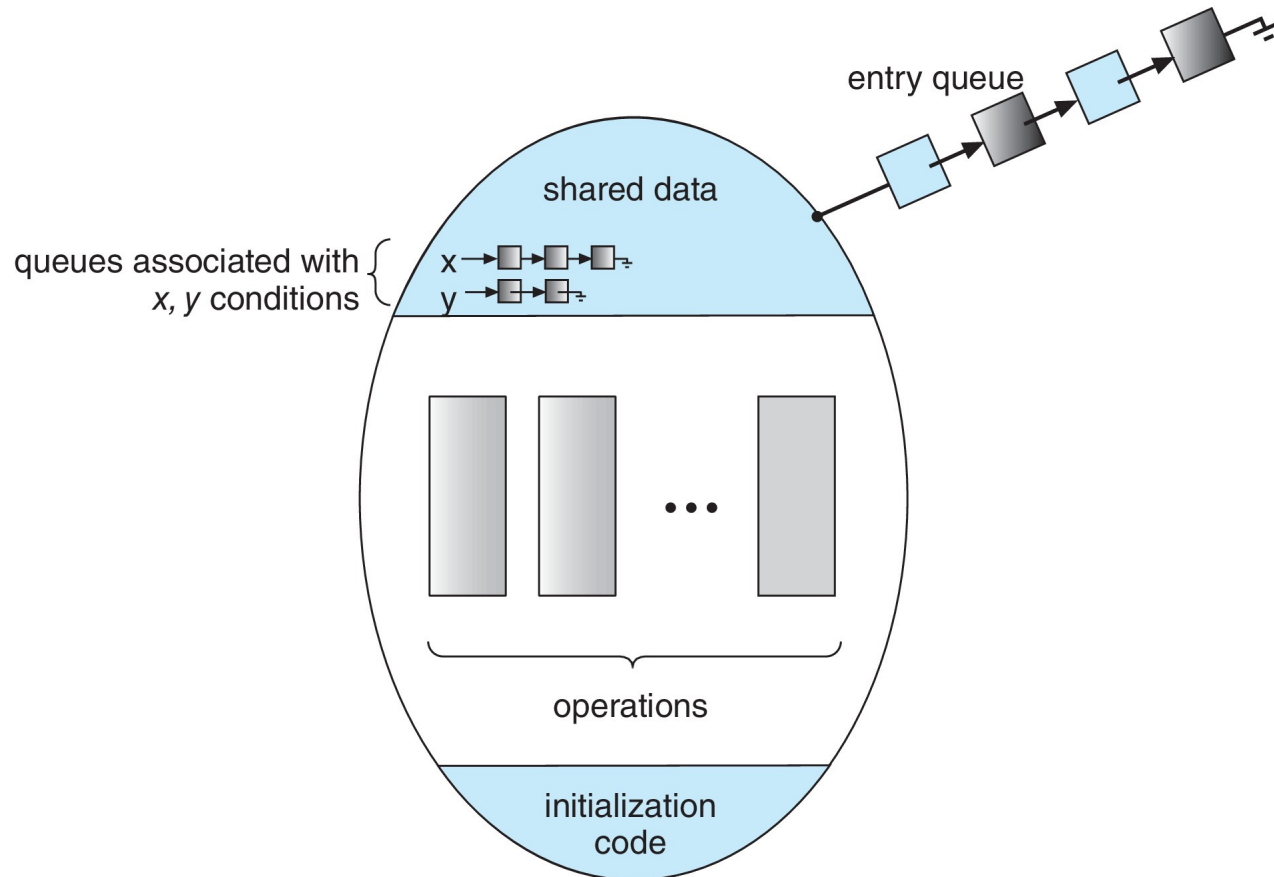
condition x, y ;

- Two operations are allowed on a condition variable:
 - **$x.\text{wait}()$** – a process that invokes the operation is suspended until **$x.\text{signal}()$**
 - **$x.\text{signal}()$** – resumes one of processes (if any) that invoked **$x.\text{wait}()$**
 - ▶ If no **$x.\text{wait}()$** on the variable, then it has no effect on the variable





Monitor with Condition Variables





Condition Variables Choices

- If process P invokes **`x.signal()`** , and process Q is suspended in **`x.wait()`** , what should happen next?
 - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
 - Both have pros and cons – language implementer can decide
 - Monitors implemented in Concurrent Pascal compromise
 - ▶ P executing signal immediately leaves the monitor, Q is resumed
 - Implemented in other languages including Mesa, C#, Java





Liveness

- Processes may have to wait indefinitely while trying to acquire a synchronization tool such as a **mutex lock** or **semaphore**.
- Waiting indefinitely violates the progress and bounded-waiting criteria discussed at the beginning of this chapter.
- **Liveness** refers to a set of properties that a system must satisfy to ensure processes make progress.
- Indefinite waiting is an example of a **liveness failure**.





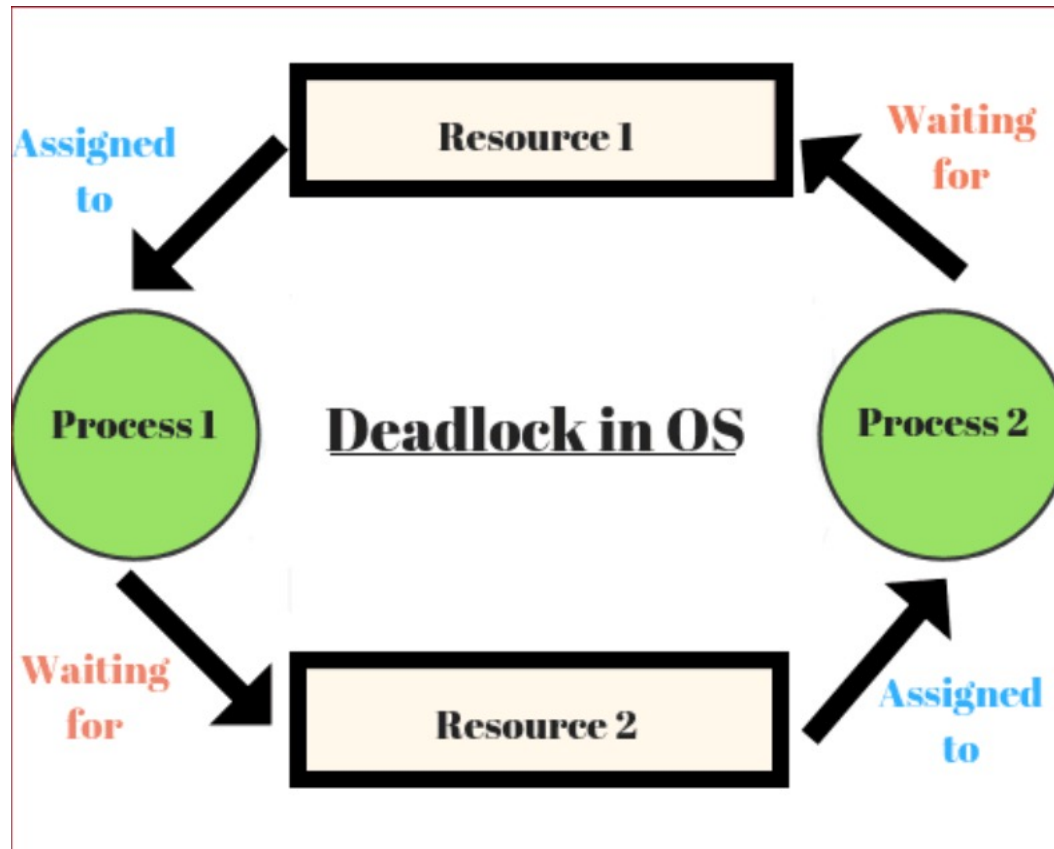
Liveness

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0	P_1
<code>wait(S) ;</code>	<code>wait(Q) ;</code>
<code>wait(Q) ;</code>	<code>wait(S) ;</code>
<code>...</code>	<code>...</code>
<code>signal(S) ;</code>	<code>signal(Q) ;</code>
<code>signal(Q) ;</code>	<code>signal(S) ;</code>

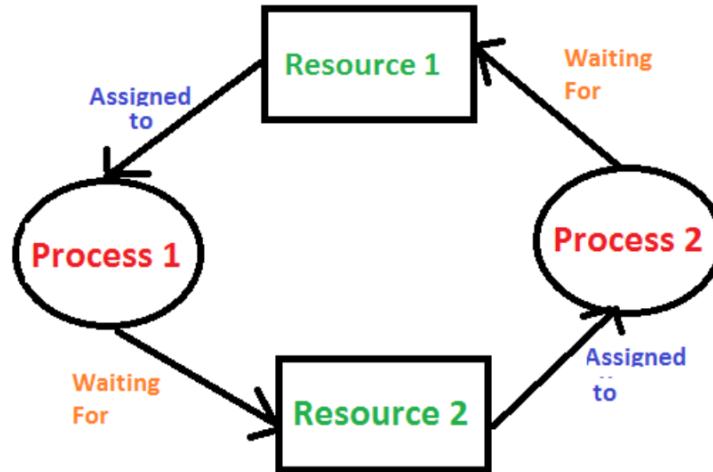
- Consider if P_0 executes `wait(S)` and P_1 `wait(Q)`. When P_0 executes `wait(Q)`, it must wait until P_1 executes `signal(Q)`
- However, P_1 is waiting until P_0 execute `signal(S)`.
- Since these `signal()` operations will never be executed, P_0 and P_1 are **deadlocked**.







Liveness



- Other forms of deadlock:
- **Starvation** – indefinite blocking
 - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Solved via **priority-inheritance protocol**

